

Baoquan Zhang

Summer Internship 2018

☎ (612) 212 9868

✉ zhan4281@umn.edu

Education

- 2015 – 2020 **Ph.D. Candidate of Computer Science**, *University of Minnesota*, Twin Cities, U.S.,
Advisor: Prof. David H.C. Du.
- Persistent Memory Systems, Key-Value Stores, RAID Systems, etc.
- 2008 – 2015 **B.E. and M.S. of Computer Science**, *Harbin Engineering University (HEU)*, China,
Co-Advisor: Prof. Jingmei Li (HEU), Prof. Dongsheng Wang (Tsinghua University).
- Distributed Data Management Systems: HDFS, Hive, Impala and etc.

Selected Research Projects

- Nov. 2017 – **NV-LSM: Log-Structured Merging on Hybrid Volatile/Non-Volatile Memory Systems**,
FAST'18 WiP, 1st Author.
- Propose the lazy-compaction including two steps – linking and merging.
 - Preliminary evaluations show that we reduce 60% of data reads/writes during the compaction.
- Dec. 2015 – **SmartRAID: a RAID-5 on Shingled Magnetic Recording (SMR) Drives**,
Sponsored by Seagate, Short Paper submitted to ATC'18, Full Paper under revision, 1st Author, HotStorage'16, Transactions on Computers (Nov. 2017), 4th Author.
- Characterize the performance of SMR drives and design a RAID-5 using alternating cleaning policy.
 - Evaluations show that SmartRAID reduces over 99% of data cleaning overheads in SMR drives.
- 2015 – 2016 **Improving Data Integrity in Linux Software RAID using Protection Information**,
Sponsored by Cray, CCGrid'18 Workshop, 1st Author.
- Implement a new software RAID module in Linux kernel compatible with T10 Protection Information.
 - Evaluation shows that new module detects data corruptions with 10% – 30% of performance overheads.
- 2013 – 2014 **Shared I/O Scheduling in Cloud for Structured Data Processing**,
BDCloud'14, 1st Author, Chinese Patent: ZL201410081163.7.
- Balance loads based on the node performance and merge queries to same data sets in a cluster.
 - Evaluations show that we improve the system performance by over 30%.

Working Experience

- Summer 2017 **Dell EMC, Eden Prairie, Minnesota**,
Software Engineering Intern, Manager: Jim Rohde,
Project: Performance improvements of the IO tracing module in storage controllers (C++).
- Conduct comprehensive evaluations on existing tracing module and identify the performance bottleneck.
 - Implement a parallel tracing module and improve performance by 50% – 400%.
- 2013 – 2014 **Research Institute of Information Technology, Tsinghua University, Beijing, China**,
Full-Time Research Assistant, Advisor: Prof. Dongsheng Wang,
Project: Construction of geo-distributed data management systems (C++, Python, Java).
- Construct a data management system using MySQL and Impala on HDFS.
 - Realize data selecting and joining in the system deployed in Beijing and Suzhou.

Miscellaneous

- | | | |
|-----------|---|--|
| 2015-2016 | ADC Graduate Fellowship | <i>University of Minnesota – Twin Cities</i> |
| 2013-2014 | National Fellowship of China (Rank 1/120) | <i>Harbin Engineering University</i> |
| 2012-2015 | 1st-Class Fellowship for Outstanding Students (3-Year) | <i>Harbin Engineering University</i> |